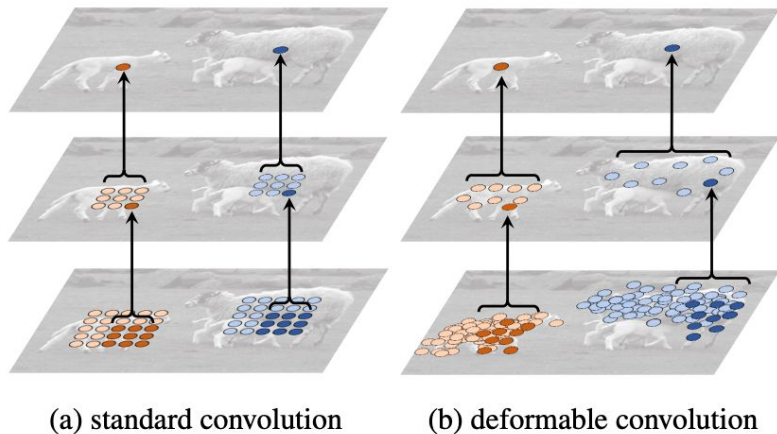


# Deformable Convolutions


PA228: Machine Learning in Image Processing

Vasilii Karmazin

# Methodology



- Can handle geometric transformations in object scale, pose, viewpoint, and part deformation
- The 2D offset vectors are learned as an addition to the regular grid sampling locations
- Modulation mechanism allow network to control signal strength from a particular location.  $\Delta m \in [0, 1]$





$$\text{Conv: } y(\mathbf{p}_0) = \sum_{\mathbf{p}_n \in \mathcal{R}} \mathbf{w}(\mathbf{p}_n) \cdot \mathbf{x}(\mathbf{p}_0 + \mathbf{p}_n) \quad (1)$$

$$\text{Def Conv: } y(\mathbf{p}_0) = \sum_{\mathbf{p}_n \in \mathcal{R}} \mathbf{w}(\mathbf{p}_n) \cdot \mathbf{x}(\mathbf{p}_0 + \mathbf{p}_n + \Delta \mathbf{p}_n) \quad (2)$$

$$\text{Modulated Def Conv: } y(p) = \sum_{k=1}^K w_k \cdot x(p + p_k + \Delta p_k) \cdot \Delta m_k \quad (3)$$


# Results

Model	Valid mIoU	Train mIoU
norm-conv	22.18%	31.03%
def-conv-v1	22.98% ( +0.80% )	32.09% ( +1.06% )
def-conv-v2	23.59% ( +1.41% )	32.92% ( +1.89% )



# Conclusion

- Deformable Convolutions works!
  - Gives +1.4% mIoU on validation
  - Adds 300-400k params, which is ~1% of the total number of parameters
  - Learning is a little slower
- My offset visualizations are not impressive ;(
- I underestimated the task:
  - ADE20k is hard dataset with 150 classes
  - ResNet50 is huge network with 33M params, which took 3-4h to train
- Always select hyperparameters from papers
- There were two major bugs in the code that forced me to retrain the models



Visualization of the offsets from the original paper